

torchgeo-bench

Two-Week Sprint Recap

~40 PRs · GPU evaluation · new models · efficiency profiling · leaderboard results

May 2026

What shipped

GPU

GPU KNN via faissknn

FAISS-backed KNN with optional CUDA path. Auto-fallback to CPU.

CLI

Typer + Rich

Replaced argparse + tqdm. Beautiful progress bars, rich tracebacks.

PROFILE

Efficiency Profiling

Throughput, GFLOPs, peak GPU mem, energy (Wh/1k), \$/inference.

MODELS

OlmoEarth + DINOv3-SAT

OlmoEarth nano/tiny/base/large. DINOv3-sat ViT-L web-pretrained.

DATA

EuroSAT Spatial Split

Geographically disjoint train/test. Harder than random split.

QUALITY

Cleanlab Label Audit

Label-quality scores across all GeoBench V1+V2 datasets.

POOL

CLS + Mean Pool Ablations

TerraTorch/ScaleMAE/Clay — pool=cls|mean|both sweep.

FIX

Silent-bug Sweep

Removed try/except covers; fixed minmax_zscore, fp16 overflow, label gaps.

GPU KNN — faissknn

PR #53

PR #55

PR #89

CPU path:

```
1 clf = KNNClassifier(n_neighbors=5, device="cpu")
2 clf.fit(x_train, y_train)
3 preds = clf.predict(x_test)
```

GPU path (opt-in via `pip install -e ".[cuda]"`):

```
1 clf = KNNClassifier(
2     n_neighbors=5,
3     device="cuda:0",
4     metric="cosine", # l2 | ip | cosine
5 )
6 clf.fit(x_train, y_train)
7 preds = clf.predict(x_test)
8 proba = clf.predict_proba(x_test)
```

Key fixes:

- `n_classes = max(y)+1` not `len(unique(y))` — avoids `IndexError` on partitions with missing class labels
- `use_fp16=False` in evaluation — raw sensor DN values (~10 000) overflow fp16 L2 distances → random KNN
- `faiss-cpu` + `faiss-cuda` namespace conflict resolved on GPU SLURM nodes

Auto-fallback:

```
1 # faissknn not installed →
2 # silently falls back to CPU faiss path
```

Typer + Rich CLI

PR #88

Before (argparse + tqdm):

```
1 parser = argparse.ArgumentParser()
2 parser.add_argument("--model", ...)
3 args = parser.parse_args()
4
5 for batch in tqdm(dataloader):
6     ...
```

After (Typer + Rich):

```
1 app = typer.Typer(no_args_is_help=True)
2
3 @app.command(context_settings={
4     "allow_extra_args": True,
5     "ignore_unknown_options": True,
6 })
7 def run(ctx: typer.Context) → None:
8     """Run benchmarks; extra args → Hydra."""
9     sys.argv = [sys.argv[0], *ctx.args]
10    hydra_main()
```

Rich progress bars:

```
1 from rich.progress import track
2
3 for batch in track(dataloader,
4                   description="Extracting"):
5     features = model(batch["image"])
```

Rich logging:

```
1 from rich.logging import RichHandler
2
3 logging.basicConfig(handlers=[
4     RichHandler(rich_tracebacks=True)
5 ])
```

Rich tables (tune_dataloader.py):

```
1 from rich.table import Table
2
3 table = Table(header_style="bold cyan")
```

Efficiency Profiling

PR #60

PR #62

PR #63

PR #67

Metrics recorded per model run:

```
1  {
2    # GPU
3    "throughput_samples_per_sec": 1420.3,
4    "peak_gpu_mem_gb":          3.2,
5    "gpu_power_w_avg":          182.0,
6    "energy_wh_per_1k_samples":  0.036,
7    "gflops":                    61.6,
8    "params_m":                  307.4,
9    # CPU
10   "throughput_samples_per_sec_cpu": 42.1,
11   # Cost
12   "cost_usd_per_1M_samples":    0.12,
13   "gco2_per_1M_samples":       18.4,
14 }
```

GFLOPs via FlopCounterMode:

```
1  with torch.profiler.FlopCounterMode() as fc:
2      model(sample)
3      gflops = fc.get_total_flops() / 1e9
```

Energy via pynvml:

```
1  import pynvml
2  pynvml.nvmlInit()
3  h = pynvml.nvmlDeviceGetHandleByIndex(0)
4  mw = pynvml.nvmlDeviceGetPowerUsage(h)
5  wh = (mw / 1000) * (elapsed_s / 3600)
```

Cost extrapolation:

```
1  # A100 spot ~$1.50/hr on Lambda
2  cost = (1_000_000 / throughput) / 3600 * 1.50
```

OlmoEarth Integration

PR #84

PR #85

Config (nano → large):

```
1 # conf/model/olmoearth_v1_base.yaml
2 _target_: torchgeo_bench.models.OlmoEarthBench
3 name: olmoearth_v1_base
4 variant: base
5 normalization: identity # model handles its own
```

Auto-rescale to S2 DN:

```
1 class OlmoEarthBench(BenchModel):
2     expected_input_unit = InputUnit.S2_DN
3
4     def _forward_patch_features(
5         self, images: Tensor, **_
6     ) → Tensor:
7         # model_native normalizer already
8         # rescaled inputs to S2 DN range
9         return self.backbone(images)
```

m-eurosat Linear Accuracy:

	Model	Score
1	OlmoEarth Large	0.976
2	OlmoEarth Base	0.975
3	DOFA Large	0.973
4	DINOv3-SAT ViT-L	0.969
5	Panopticon	0.968

Dominates S2 datasets — EuroSAT, BigEarthNet, So2Sat.
Nano/Tiny competitive at a fraction of the size.

CLS Token + Pool Ablations

[PR #73](#)[PR #74](#)[PR #75](#)[PR #76](#)

`pool` kwarg across TerraTorch wrappers:

```
1 class TerramindBench(BenchModel):
2     def _forward_patch_features(
3         self, images: Tensor, **_
4     ) → Tensor:
5         out = self.model.encode(images)
6         if self.pool == "cls":
7             return out[:, 0]
8         elif self.pool == "mean":
9             return out[:, 1:].mean(1)
10        else: # "both"
11            return torch.cat([
12                out[:, 0],
13                out[:, 1:].mean(1),
14            ], dim=1)
```

Config variants:

```
1 name: tt_clay_v1_5_base_cls # pool: cls
```

Finding: CLS helps on ForestNet / Brick Kiln; patch-mean wins on EuroSAT / BigEarthNet.

Terramind has no CLS token — `_cls` configs dropped (#76).

m-brick-kiln Linear Accuracy:

Model		Score
1	DINOv3-SAT ViT-L	0.976
2	Clay v1.5 Base+CLS	0.975
3	DOFA Base	0.974
4	DOFA Large	0.974

EuroSAT Spatial Split + Cleanlab

PR #50

PR #52

Geographically disjoint train/test:

```
1 # Standard: random split
2 # → models exploit spatial autocorrelation
3
4 # Spatial: tiles from different regions
5 # → true generalization test
6 ds = EuroSATSpatialBench(
7     root=DATA_ROOT,
8     partition="default",
9     bands=["B02", "B03", "B04", "B08"],
10 )
```

OlmoEarth Large KNN: **0.959** vs 0.956 random split

OlmoEarth Base Linear: **0.978** vs 0.975 random split

Cleanlab label audit:

```
1 from cleanlab.filter import find_label_issues
2
3 issues = find_label_issues(
4     labels=y_train,
5     pred_probs=pred_proba,
6     return_indices_ranked_by=
7         "self_confidence",
8 )
9 # → flags likely mislabeled samples
```

Applied across all GeoBench V1+V2. Results in

`results/cleanlab/`.

Surfaces annotation noise in BigEarthNet, ForestNet, TreeSatAI — useful for reweighting or curriculum training.

Leaderboard — EuroSAT

RESULTS

m-eurosat (Accuracy)

	Model	KNN	Lin
1	OlmoEarth Large	.956	.976
2	OlmoEarth Base	.946	.975
3	Panopticon	.948	.968
4	DOFA Large	.936	.973
5	DINOv3-SAT ViT-L	—	.969

eurosat-spatial (Accuracy)

	Model	KNN	Lin
1	OlmoEarth Large	.959	.977
2	OlmoEarth Base	.941	.978
3	ResNet50-S2 MoCo	.933	—
4	DINOv3 ViT-L	.932	—
5	Panopticon	.930	.962

Leaderboard — BigEarthNet

RESULTS

m-bigearthnet (micro-mAP)

	Model	KNN	Lin
1	OlmoEarth Large	.664	.764
2	OlmoEarth Base	.658	.769
3	Panopticon	.652	—
4	Terramind Large	.628	.750
5	DINOv3-SAT ViT-L	.625	.741

benv2 (micro-mAP)

	Model	KNN	Lin
1	OlmoEarth Base	.735	.853
2	OlmoEarth Large	.728	.850
3	OlmoEarth Tiny	.716	—
4	Panopticon	.716	—
5	Terramind Large	.712	.846

Leaderboard — ForestNet + So2Sat

RESULTS

m-forestnet (Accuracy)

	Model	KNN	Lin
1	Panopticon	.427	—
2	DINOv3-SAT ViT-L	.425	.582
3	Clay v1.5 Base+CLS	.414	.556
4	ScaleMAE Large+CLS	.403	.569
5	Clay v1.5 Base	.395	—

m-so2sat (Accuracy)

	Model	KNN	Lin
1	OlmoEarth Base	.576	.712
2	OlmoEarth Large	.560	—
3	Panopticon	.532	—
4	OlmoEarth Tiny	.506	—
5	DINOv3 ViT-L	.490	—

Linear leader: Terramind Base .739

Leaderboard — Brick Kiln + PV4GER + TreeSatAI

RESULTS

m-brick-kiln (Acc)

Model		KNN
1	DOFA Large	.969
2	DOFA Base	.965
3	Clay Base	.964
4	Clay Base+CLS	.960
5	DINOv3-SAT ViT-L	—

Lin leader: DINOv3-SAT .976

m-pv4ger (Acc)

Model		KNN
1	DOFA Large	.967
2	DOFA Base	.965
3	DINOv3 ViT-L	.964
4	Panopticon	.958
5	ScaleMAE+CLS	.956

Lin leader: DINOv3 ViT-L .974

treesatai (mAP)

Model		KNN
1	DINOv3-SAT ViT-L	.477
2	Clay Base+CLS	.475
3	DOFA Large	.474
4	DINOv3 ViT-L	.473
5	OlmoEarth Nano	.470

Lin leader: DINOv3-SAT .682

Efficiency — Throughput vs Accuracy

PR #60-#80

GPU throughput (img/s) — m-eurosat

Model	img/s	GFLOPs	Acc
ResNet-50 MoCo	3 193	8	.76
DOFA Base	1 747	36	.71
Terramind Base	1 680	36	.68
OlmoEarth Tiny	780	9	.73
OlmoEarth Nano	789	2	.67
DOFA Large	581	124	.73
DINOv3-SAT ViT-L	346	165	.74

Pareto surprises:

WINNER **ResNet-50 MoCo** — 3 193 img/s, 8 GFLOPs, 23M params, yet **.76 avg acc.** Matches or beats every ViT-scale model while running 5–10× faster.

EFFICIENT **OlmoEarth Nano** — 3.6M params, 0.6 GB peak VRAM, 1.6 GFLOPs. Matches Panopticon accuracy at 5× lower cost.

EFFICIENT **OlmoEarth Tiny** (14M) ties DOFA Large (337M) and DINOv3-SAT (304M) at 2× the throughput.

Worst value: Prithvi v1/v2 — 1 700 img/s, 35 GFLOPs, only .56–.60 acc. Same throughput tier as DOFA Base but far lower accuracy.

Intrinsic Dimensionality Analysis

METHOD: INTRINSIC_DIM

id_TwoNN (avg across datasets)

Model	id_TwoNN	Feat dim	Avg acc
EarthLoc ResNet-50	44	4 096	.59
OlmoEarth Large	22	1 024	.74
DOFA Large	18	1 024	.73
DINOv3-SAT ViT-L	16	1 024	.74
ResNet-50 MoCo	15	2 048	.76
OlmoEarth Nano	14	128	.67
Panopticon	13	768	.66

Pattern: High intrinsic dim ↔ high accuracy. Top models (OlmoEarth, DOFA, DINOv3-SAT, ResNet-50 MoCo) all ID > 14.

Outliers:

TASK GAP **EarthLoc ResNet-50** — id = 44 (highest), but only .59 acc. Trained for geo-localization, not classification. High dimensionality without discriminative structure.

SURPRISE **OlmoEarth Nano** — 128-d output yet id_TwoNN = 14, matching DINOv3-SAT ViT-L (1024-d). Packs equivalent intrinsic complexity into 8× fewer dimensions.

FLAT **Prithvi cluster** — id ≈ 7, just above imagestats (6) and RCF (5). Geometrically flat despite 768–1280-d features — explains consistently low accuracy.

Key Findings

S2 multispectral (EuroSAT, BigEarthNet, So2Sat)

→ **OlmoEarth** dominates KNN + Linear

RGB / aerial (ForestNet, PV4GER, Brick Kiln)

→ **DINOv3-SAT** or **Panopticon** lead

Multi-label (BigEarthNet, BENV2, TreeSatAI)

→ OlmoEarth top KNN; Terramind competitive linear

✗ Terramind **Base** beats Terramind **Large** on So2Sat linear (0.739 vs 0.712)

✗ CLS token **collapses** on EuroSAT for Prithvi/Clay — patch-mean wins

~40 PRs · GPU KNN · Typer+Rich · efficiency profiling
· OlmoEarth · spatial splits · cleanlab

"OlmoEarth leads on S2 · DINOv3-SAT leads aerial · Terramind surprises on So2Sat"